

B.E CSE VI Q BATCH
CS6301-MACHINE LEARNING

**PREDICTIVE MAINTENANCE OF AIRCRAFT ENGINE
USING LSTM NETWORKS**

TEAM NUMBER : 03

TEAM MEMBERS:

VIGNESH G - 2018103078

DEEKSHITH M -2019103014

AJITESH M - 2019103503

HEMANTH D - 2019103020

PROBLEM STATEMENT

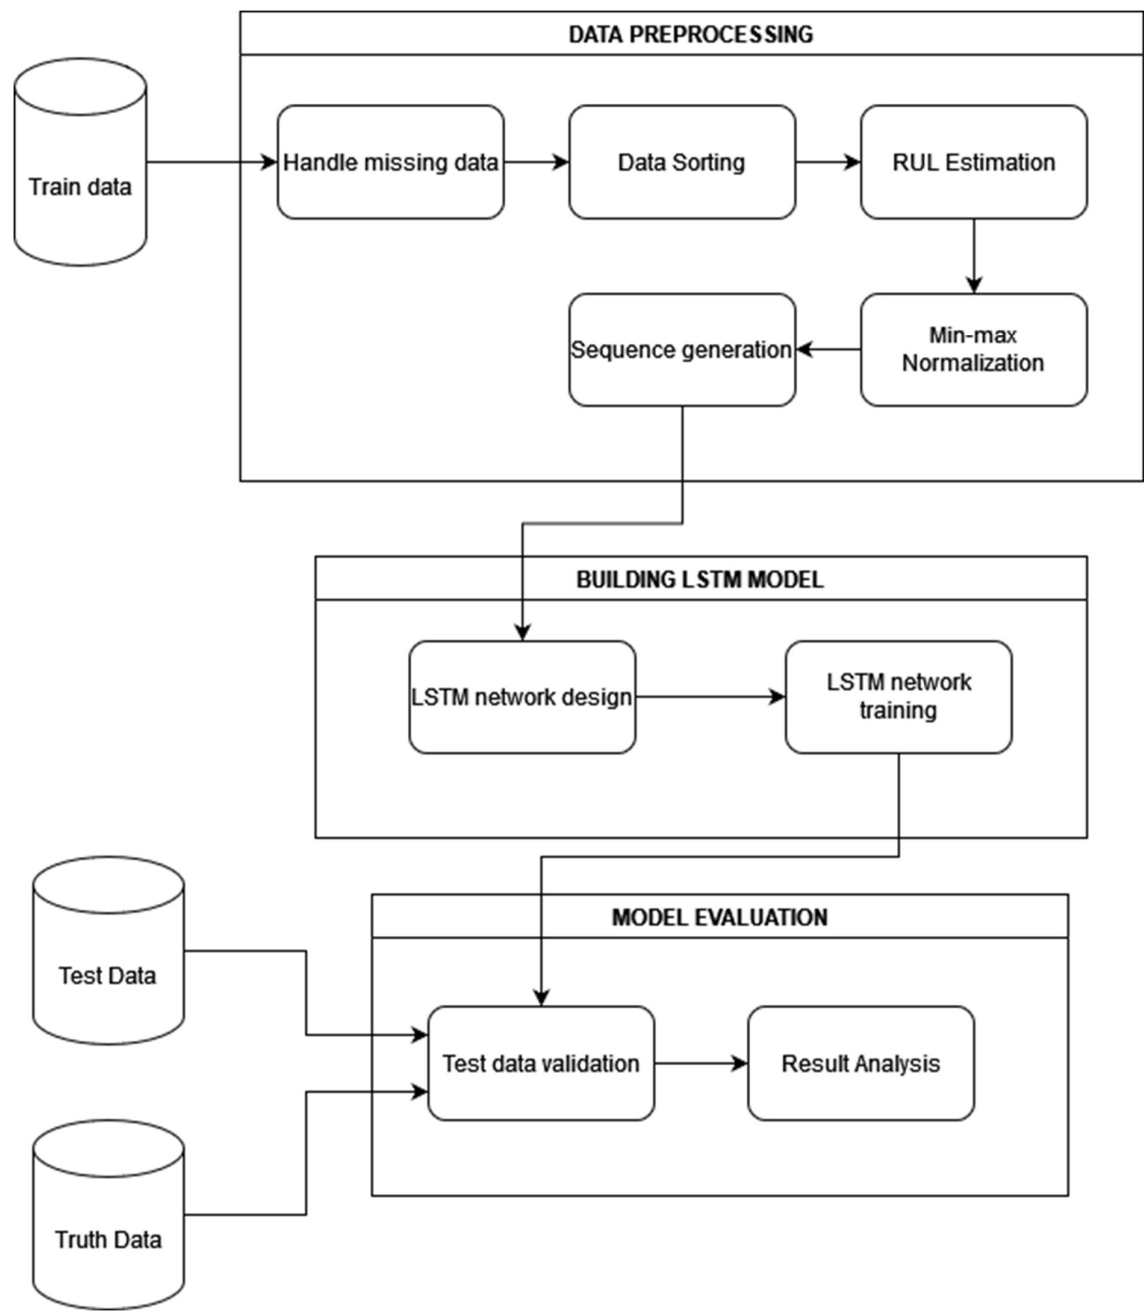
What if a part of aircraft could let know when the aircraft component needed to be replaced or repaired? It can be done with continuous data collection, monitoring and advanced analytics. In the aviation industry, predictive maintenance promises increased reliability as well as improved supply chain and operational performance. The main goal is to ensure that the engines work correctly under all conditions and there is no risk of failure. The main source of data regarding the health of the engine is measured during the real time. Several variables are calculated, including fan speed, quantity and oil pressure and environmental variables such as temperature, air speed.

Aspects related to the maintenance have become especially failure of a have catastrophic consequences. Current systems have the ability to warn inform only when the failure occurred. An early warning system that predicts the occurrence of component failure is required. A machine learning approach, using LSTM networks, is proposed to predict the RUL (Remaining Useful Life) by analysing failure patterns in the past. Training of LSTM networks are carried out on a high performance large-scale processing engine.

OBJECTIVE

A machine learning approach through the use of Long Short Term Memory (LSTM) networks is used to analyse sensor time series sequences to estimate the RUL of turbofan engines, and we provide an analysis where we show how the LSTM performance changes when varying its internal hyperparameters.

OVERALL ARCHITECTURE DIAGRAM



DETAILS OF MODULE DESIGN

- 1.Data pre-processing
2. Building LSTM network
3. Training the model
4. Testing the model and Result analysis.

1.DATA PRE-PROCESSING:

The dataset has inputs.

- Training data: It is the aircraft engine run-to-failure data.
- Testing data: It is the aircraft engine operating data without failure events recorded.
- Ground truth data: It contains the information of true remaining cycles for each engine in the testing data.

The training data ("train_FD001.txt") consists of multiple multivariate time series with "cycle" as the time unit, together with 21 sensor readings for each cycle. Each time series can be assumed as being generated from a different engine of the same type. Each engine is assumed to start with different degrees of initial wear and manufacturing variation, and this information is unknown to the user. In this simulated data, the engine is assumed to be operating normally at the start of each time series. It starts to degrade at some point during the series of the operating cycles. The degradation progresses and grows in magnitude. When a predefined threshold is reached, then the engine is considered unsafe for further operation. In other words, the last cycle in each time series can be considered as the failure point of the corresponding engine. Taking the sample training data shown in the following table as an example, the engine with id=1 fails at cycle 192, and engine with id=2 fails at cycle 287. with this information we can calculate RUL for training dataset

The testing data ("test_FD001.txt") has the same data schema as the training data. The only difference is that the data does not indicate when the failure occurs (in other words, the last time period does NOT represent the failure point). Taking the sample testing data shown in the following table as an example, the engine with id=1 runs from cycle 1 through cycle 31. It is not shown how many more cycles this engine can last before it fails.

The ground truth data ("RUL_FD001.txt") provides the number of remaining working cycles for the engines in the testing data. Taking the sample ground truth data shown in the following table as an example, the engine with id=1 in the testing data can run another 112 cycles before it fails.

2.BUILDING THE MODEL

Long short term Memory Network Known as LSTM. It is one of the best kind of RNN with capability of avoiding gradient dispersion. It is designed to avoid long term dependencies. LSTM Cells are where data is transfers and updated, cell states of the LSTM is changed as compared to RNN. Network is based on short term states, long term states and its three gates: input gate, output gate and forget gate. Where f_t is forget gate use to forget the information that is no longer required.




Figure 2: LSTM architecture

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

here σ is activation function x is input to the gate and b is bias vector.

The input gate contains two path one for the new input and second for vector I generated by forget gate, which is use to modify the cell state.

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$C_t^{\sim} = \tanh(W_c \cdot [h_{t-1}, x_t] + b_c)$$

Where W are the weight matrices and b are the bias vectors for input gate with activation function of \tanh . And updated state of gate is:

$$C_t = f_t \otimes C_{t-1} + i_t \otimes C_t^{\sim}$$

output gates have two parts as input gate.

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t \otimes \tanh(C_t)$$

3.TRAINING THE MODEL.

We input our training dataset for which we know the Remaining Useful lifetime[RUL], our LSTM model takes this as input and calculates the RUL. From our experience, we noticed that the obtained results can be very different when varying the model

architecture so the base idea of such an analysis is to show how the model performance, in terms of Root Mean Squared Error (RMSE), changes over a different set of hyperparameters. RMSE is commonly used in supervised learning applications, as RMSE uses and needs true measurements at each predicted data point.

Root mean square error can be expressed as

$$RMSE = \sqrt{\frac{\sum_{i=1}^N \|y(i) - \hat{y}(i)\|^2}{N}},$$

where N is the number of data points, y(i) is the i-th measurement, and $\hat{y}(i)$ is its corresponding prediction.

We were able to analyze the RMSE variations when changing the hyperparameters values and understand which are the elements that affects it mostly. The result of such an analysis is very important because it allows to have a better understanding on how this kind of models work and how we can improve them by changing their configuration. Finally we fix the hyperparameters which give us high accuracy.

4. TESTING THE MODEL AND RESULT ANALYSIS.

We finally apply the Training dataset whose ground truth Remaining useful lifetime is given in a separate dataset into our trained and optimized LSTM model, we obtain the output for each case of aircraft engine based on the aircraft id. The RUL predicted and actual ground truth is compared to analyze the result.

Model performance on the test data showing the differences between the predictions made and the ground truth values where the generic i-th difference has been computed as:

$$d_i = \hat{y}_i - y_i$$

In this sense, a value higher than zero means that our model made an optimistic prediction while a value lower than zero means a pessimistic prediction. Nevertheless, the majority of them presents a difference near to zero meaning that the model is able to correctly predict the RUL with a very good level of accuracy. We analyze the accuracy, error rate using various graph plotting packages such as pyplot, seaborn to get a conclusion of how our LSTM model performs for predicting Remaining useful lifetime for Training dataset.

IMPLEMENTATION

MODULE-1:DATA PREPROCESSING

Required packages and libraries needed for building the model are imported.

```
1 import keras
2 import keras.backend as K
3 from keras.layers.core import Activation
4 from keras.models import Sequential,load_model
5 from keras.layers import Dense, Dropout, LSTM
6
7 import pandas as pd
8 import numpy as np
9 import matplotlib.pyplot as plt
10 import os
11 from sklearn import preprocessing
```

```
[3] 1 # Setting seed for reproducibility
      2 np.random.seed(1234)
      3 PYTHONHASHSEED = 0
```

The dataset containing training and testing data are uploaded and the null values are removed in order to avoid false information.

```
1 # read training data - It is the aircraft engine run-to-failure data.
2 # read test data - It is the aircraft engine operating data without failure events recorded.
3 # read ground truth data - It contains the information of true remaining cycles for each
4 # engine in the testing data.
5 train_df = pd.read_csv('PM_train.txt', sep=" ", header=None)
6 test_df = pd.read_csv('PM_test.txt', sep=" ", header=None)
7 truth_df = pd.read_csv('PM_truth.txt', sep=" ", header=None)
```

```
[5] 1 # Drop missing data columns(redundant)
      2 train_df.drop(train_df.columns[[26, 27]], axis=1, inplace=True)
      3 test_df.drop(test_df.columns[[26, 27]], axis=1, inplace=True)
      4 truth_df.drop(truth_df.columns[[1]], axis=1, inplace=True)
```

```
[6] 1 # Sorting and indicating columns
      2 train_df.columns = ['id', 'cycle', 'setting1', 'setting2', 'setting3', 's1', 's2', 's3',
      3 |                  's4', 's5', 's6', 's7', 's8', 's9', 's10', 's11', 's12', 's13', 's14',
      4 |                  's15', 's16', 's17', 's18', 's19', 's20', 's21']
      5
      6 train_df = train_df.sort_values(['id','cycle'])
      7
      8 test_df.columns = ['id', 'cycle', 'setting1', 'setting2', 'setting3', 's1', 's2', 's3',
      9 |                  's4', 's5', 's6', 's7', 's8', 's9', 's10', 's11', 's12', 's13', 's14',
     10 |                  's15', 's16', 's17', 's18', 's19', 's20', 's21']
```

```
[7] 1 train_df
```

	id	cycle	setting1	setting2	setting3	s1	s2	s3	s4	s5	...	s12	s13	s14	s15	s16	s17	s18	s19	s20	s21
0	1	1	-0.0007	-0.0004	100.0	518.67	641.82	1589.70	1400.60	14.62	...	521.66	2388.02	8138.62	8.4195	0.03	392	2388	100.0	39.06	23.4190
1	1	2	0.0019	-0.0003	100.0	518.67	642.15	1591.82	1403.14	14.62	...	522.28	2388.07	8131.49	8.4318	0.03	392	2388	100.0	39.00	23.4236
2	1	3	-0.0043	0.0003	100.0	518.67	642.35	1587.99	1404.20	14.62	...	522.42	2388.03	8133.23	8.4178	0.03	390	2388	100.0	38.95	23.3442
3	1	4	0.0007	0.0000	100.0	518.67	642.35	1582.79	1401.87	14.62	...	522.86	2388.08	8133.83	8.3682	0.03	392	2388	100.0	38.88	23.3739
4	1	5	-0.0019	-0.0002	100.0	518.67	642.37	1582.85	1406.22	14.62	...	522.19	2388.04	8133.80	8.4294	0.03	393	2388	100.0	38.90	23.4044
...
20626	100	196	-0.0004	-0.0003	100.0	518.67	643.49	1597.98	1428.63	14.62	...	519.49	2388.26	8137.60	8.4956	0.03	397	2388	100.0	38.49	22.9735
20627	100	197	-0.0016	-0.0005	100.0	518.67	643.54	1604.50	1433.58	14.62	...	519.68	2388.22	8136.50	8.5139	0.03	395	2388	100.0	38.30	23.1594
20628	100	198	0.0004	0.0000	100.0	518.67	643.42	1602.46	1428.18	14.62	...	520.01	2388.24	8141.05	8.5646	0.03	398	2388	100.0	38.44	22.9333
20629	100	199	-0.0011	0.0003	100.0	518.67	643.23	1605.26	1426.53	14.62	...	519.67	2388.23	8139.29	8.5389	0.03	395	2388	100.0	38.29	23.0640
20630	100	200	-0.0032	-0.0005	100.0	518.67	643.85	1600.38	1432.14	14.62	...	519.30	2388.26	8137.33	8.5036	0.03	396	2388	100.0	38.37	23.0522

20631 rows x 26 columns

Then the Remaining Useful Life is calculated using the number of cycles given in the dataset for training dataset.

```
[8] 1 # Data Preprocessing - Train data
2 # Data Labeling - generate column RUL(Remaining Usefull Life or Time to Failure)
3 rul = pd.DataFrame(train_df.groupby('id')['cycle'].max()).reset_index()
4 rul.columns = ['id', 'max']
5 train_df = train_df.merge(rul, on='id', how='left')
6 train_df['RUL'] = train_df['max'] - train_df['cycle']
7 train_df.drop('max', axis=1, inplace=True)
8
9 # MinMax normalization (from 0 to 1)
10 train_df['cycle_norm'] = train_df['cycle']
11 cols_normalize = train_df.columns.difference(['id', 'cycle', 'RUL', 'label1', 'label2'])
12 min_max_scaler = preprocessing.MinMaxScaler()
13 norm_train_df = pd.DataFrame(min_max_scaler.fit_transform(train_df[cols_normalize]),
14                               columns=cols_normalize,
15                               index=train_df.index)
16 join_df = train_df[train_df.columns.difference(cols_normalize)].join(norm_train_df)
17 train_df = join_df.reindex(columns=train_df.columns)
18
19 print(train_df)
```

	id	cycle	setting1	setting2	setting3	s1	s2	s3	...	s12	s13	s14	s15	s16	s17	s18	s19	...	s20	s21
0	1	1	0.459770	0.166667	0.0	0.0	0.183735	0.406802	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
1	1	2	0.609195	0.250000	0.0	0.0	0.283133	0.453019	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
2	1	3	0.252874	0.750000	0.0	0.0	0.343373	0.369523	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
3	1	4	0.540230	0.500000	0.0	0.0	0.343373	0.256159	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
4	1	5	0.390805	0.333333	0.0	0.0	0.349398	0.257467	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
...
20626	100	196	0.477011	0.250000	0.0	0.0	0.686747	0.587312	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
20627	100	197	0.408046	0.083333	0.0	0.0	0.701807	0.729453	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
20628	100	198	0.522989	0.500000	0.0	0.0	0.665663	0.684979	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
20629	100	199	0.436782	0.750000	0.0	0.0	0.608434	0.746021	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
20630	100	200	0.316092	0.083333	0.0	0.0	0.795181	0.639634	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
...
0			0.309757	0.0	...	0.199608	0.363986	0.0	0.333333	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
1			0.352633	0.0	...	0.162813	0.411312	0.0	0.333333	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
2			0.370527	0.0	...	0.171793	0.357445	0.0	0.166667	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
3			0.331195	0.0	...	0.174889	0.166603	0.0	0.333333	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
4			0.404625	0.0	...	0.174734	0.402078	0.0	0.416667	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
...
20626			0.782917	0.0	...	0.194344	0.656791	0.0	0.750000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
20627			0.866475	0.0	...	0.188668	0.727203	0.0	0.583333	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
20628			0.775321	0.0	...	0.212148	0.922278	0.0	0.833333	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
20629			0.747468	0.0	...	0.203065	0.823394	0.0	0.583333	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
20630			0.842167	0.0	...	0.192951	0.687572	0.0	0.666667	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
...
0			0.713178	0.724662	191	0.000000														
1			0.666667	0.731014	190	0.002770														
2			0.627907	0.621375	189	0.005540														
3			0.573643	0.662386	188	0.008310														
4			0.589147	0.704502	187	0.011080														
...														
20626			0.271318	0.109500	4	0.540166														
20627			0.124031	0.366197	3	0.542936														
20628			0.232558	0.053991	2	0.545706														
20629			0.116279	0.234466	1	0.548476														
20630			0.178295	0.218172	0	0.551247														

[20631 rows x 28 columns]

Data values are normalized in order to make data uniform across various engines with different RUL range.

```

1 # Data Preprocessing - Test data
2 # MinMax normalization (from 0 to 1)
3 test_df['cycle_norm'] = test_df['cycle']
4 norm_test_df = pd.DataFrame(min_max_scaler.transform(test_df[cols_normalize]),
5                             columns=cols_normalize,
6                             index=test_df.index)
7 test_join_df = test_df[test_df.columns.difference(cols_normalize)].join(norm_test_df)
8 test_df = test_join_df.reindex(columns=test_df.columns)
9 test_df = test_df.reset_index(drop=True)
10
11 # We use the ground truth dataset to generate labels for the test data.
12 # generate column max for test data
13 rul = pd.DataFrame(test_df.groupby('id')['cycle'].max()).reset_index()
14 rul.columns = ['id', 'max']
15 truth_df.columns = ['more']
16 truth_df['id'] = truth_df.index + 1
17 truth_df['max'] = rul['max'] + truth_df['more']
18 truth_df.drop('more', axis=1, inplace=True)
19
20 # generate RUL for test data
21 test_df = test_df.merge(truth_df, on=['id'], how='left')
22 test_df['RUL'] = test_df['max'] - test_df['cycle']
23 test_df.drop('max', axis=1, inplace=True)
24
25 print(test_df)

```

```

0      id  cycle  setting1  setting2  setting3  s1      s2      s3  \
0      1      1  0.632184  0.750000      0.0  0.0  0.545181  0.310661
1      1      2  0.344828  0.250000      0.0  0.0  0.150602  0.379551
2      1      3  0.517241  0.583333      0.0  0.0  0.376506  0.346632
3      1      4  0.741379  0.500000      0.0  0.0  0.370482  0.285154
4      1      5  0.580460  0.500000      0.0  0.0  0.391566  0.352082
...
13091 100    194  0.781609  0.500000      0.0  0.0  0.611446  0.619359
13092 100    195  0.436782  0.416667      0.0  0.0  0.605422  0.537388
13093 100    196  0.465517  0.250000      0.0  0.0  0.671687  0.482014
13094 100    197  0.281609  0.583333      0.0  0.0  0.617470  0.522128
13095 100    198  0.574713  0.750000      0.0  0.0  0.524096  0.666667

      s4  s5  ...  s14  s15  s16      s17  s18  s19  \
0  0.269413  0.0  ...  0.132160  0.308965  0.0  0.333333  0.0  0.0
1  0.222316  0.0  ...  0.204768  0.213159  0.0  0.416667  0.0  0.0
2  0.322248  0.0  ...  0.155640  0.458638  0.0  0.416667  0.0  0.0
3  0.408001  0.0  ...  0.170090  0.257022  0.0  0.250000  0.0  0.0
4  0.332039  0.0  ...  0.152751  0.300885  0.0  0.166667  0.0  0.0
...
13091 0.566172  0.0  ...  0.584890  0.564063  0.0  0.500000  0.0  0.0
13092 0.671843  0.0  ...  0.572350  0.485956  0.0  0.583333  0.0  0.0
13093 0.414754  0.0  ...  0.605326  0.507888  0.0  0.583333  0.0  0.0
13094 0.626435  0.0  ...  0.622046  0.562524  0.0  0.583333  0.0  0.0
13095 0.721472  0.0  ...  0.591908  0.636399  0.0  0.666667  0.0  0.0

      s20  s21  cycle_norm  RUL
0  0.558140  0.661834  0.000000  142
1  0.682171  0.686827  0.002770  141
2  0.728682  0.721348  0.005540  140
3  0.666667  0.662110  0.008310  139
4  0.658915  0.716377  0.011080  138
...
13091 0.395349  0.418669  0.534626  24
13092 0.333333  0.528721  0.537396  23
13093 0.372093  0.429301  0.540166  22
13094 0.403101  0.518779  0.542936  21
13095 0.434109  0.402237  0.545706  20

[13096 rows x 28 columns]

```

Data is reshaped into a sequence of length 60. Each sequence contains samples, time steps, features. Engine whose values are less than sequence length are omitted, because it cannot be used for prediction. LSTM model takes sequence of input and find RUL for first value of next sequence.

```
1 # Window size extension to 60
2 sequence_length = 60
3
4 # function to reshape features into (samples, time steps, features)
5 def gen_sequence(id_df, seq_length, seq_cols):
6     data_matrix = id_df[seq_cols].values
7     num_elements = data_matrix.shape[0]
8     for start, stop in zip(range(0, num_elements - seq_length), range(seq_length, num_elements)):
9         yield data_matrix[start:stop, :]
10
11
12 # pick the feature columns
13 sensor_cols = ['s' + str(i) for i in range(1, 22)]
14 sequence_cols = ['setting1', 'setting2', 'setting3', 'cycle_norm']
15 sequence_cols.extend(sensor_cols)
16
17 # print(sequence_cols)
18 # val is a list of 192 - 60 = 142 bi-dimensional array (60 rows x 25 columns)
19 val = list(gen_sequence(train_df[train_df['id'] == 1], sequence_length, sequence_cols))
20 print(len(val))
21
22 # generator for the sequences
23 # transform each id of the train dataset in a sequence
24 seq_gen = (list(gen_sequence(train_df[train_df['id'] == id], sequence_length, sequence_cols))
25            for id in train_df['id'].unique())
26
27 # generate sequences and convert to numpy array
28 seq_array = np.concatenate(list(seq_gen)).astype(np.float32)
29 print(seq_array.shape)
30
```

```
1 # function to generate labels
2 def gen_labels(id_df, seq_length, label):
3     data_matrix = id_df[label].values
4     num_elements = data_matrix.shape[0]
5     return data_matrix[seq_length:num_elements, :]
6
7 # generate labels
8 label_gen = [gen_labels(train_df[train_df['id'] == id], sequence_length, ['RUL'])
9              for id in train_df['id'].unique()]
10
11 label_array = np.concatenate(label_gen).astype(np.float32)
12 print(label_array.shape)
13 print(label_array[:10])
14
```

```
(14631, 1)
[[131.]
 [130.]
 [129.]
 [128.]
 [127.]
 [126.]
 [125.]
 [124.]
 [123.]
 [122.]]
```

MODULE-2: BUILDING LSTM MODEL

2 layers of LSTM cells are built, with 100 units in the first layer and 50 units in the second layer. In the final layer, sigmoid activation function is included.

```
1 # Modeling
2 model_path = 'regression_model.h5'
3
4 def r2_keras(y_true, y_pred):
5     SS_res = K.sum(K.square(y_true - y_pred))
6     SS_tot = K.sum(K.square(y_true - K.mean(y_true)))
7     return (1 - SS_res / (SS_tot + K.epsilon()))
8
9 # Network Architecture
10 # The first layer is an LSTM layer with 100 units followed by another LSTM layer with 60 units.
11 # Dropout is also applied after each LSTM layer to control overfitting.
12 # Final layer is a Dense output layer with single unit and linear activation
13 # since this is a regression problem.
14 nb_features = seq_array.shape[2]
15 nb_out = label_array.shape[1]
16
17 try:
18     f = open(model_path)
19     print("Trained model already exists")
20
21 except IOError:
22     print("Initialize a model")
23     model = Sequential()
24     model.add(LSTM(
25         input_shape=(sequence_length, nb_features),
26         units=100,
27         return_sequences=True))
28     model.add(Dropout(0.3))
29     model.add(LSTM(
30         units=sequence_length,
31         return_sequences=False))
32     model.add(Dropout(0.3))
33     model.add(Dense(units=nb_out))
34     model.add(Activation("linear"))
35     model.compile(loss='mean_squared_error', optimizer='rmsprop', metrics=['mae', r2_keras])
36
37     print(model.summary())
```

Initialize a model
Model: "sequential"


Layer (type)	Output Shape	Param #
=====		
lstm (LSTM)	(None, 60, 100)	50400
dropout (Dropout)	(None, 60, 100)	0
lstm_1 (LSTM)	(None, 60)	38640
dropout_1 (Dropout)	(None, 60)	0
dense (Dense)	(None, 1)	61
activation (Activation)	(None, 1)	0
=====		
Total params: 89,101		
Trainable params: 89,101		
Non-trainable params: 0		

MODULE-3: TRAINING LSTM MODEL


The model takes the dataset as the input and calculates the Remaining Useful Lifetime. The loss is calculated after each epoch.

```
38
39 bs = 400
40 # fit the network
41 history = model.fit(seq_array, label_array, epochs=100, batch_size=bs, validation_split=0.1, verbose=1,
42                    callbacks=[keras.callbacks.EarlyStopping(monitor='val_loss', min_delta=0, patience=10, verbose=0, mode='min'),
43                               keras.callbacks.ModelCheckpoint(model_path, monitor='val_loss', save_best_only=False,
44                                                                mode='min', verbose=0)])
45
46
47 # list all data in history
48 print(history.history.keys())
```

```
1 # summarize history for R^2
2 fig_acc = plt.figure(figsize=(10, 10))
3 plt.plot(history.history['r2_keras'])
4 plt.plot(history.history['val_r2_keras'])
5 plt.title('model r^2')
6 plt.ylabel('R^2')
7 plt.xlabel('epoch')
8 plt.legend(['train', 'test'], loc='upper left')
9 plt.show()
10 fig_acc.savefig("model_r2.png")
11
```




```
12 # summarize history for MAE
13 fig_acc = plt.figure(figsize=(10, 10))
14 plt.plot(history.history['mae'])
15 plt.plot(history.history['val_mae'])
16 plt.title('model MAE')
17 plt.ylabel('MAE')
18 plt.xlabel('epoch')
19 plt.legend(['train', 'test'], loc='upper left')
20 plt.show()
21 fig_acc.savefig("model_mae.png")
22
```



```

23 # summarize history for Loss
24 fig_acc = plt.figure(figsize=(10, 10))
25 plt.plot(history.history['loss'])
26 plt.plot(history.history['val_loss'])
27 plt.title('model loss')
28 plt.ylabel('loss')
29 plt.xlabel('epoch')
30 plt.legend(['train', 'test'], loc='upper left')
31 plt.show()
32 fig_acc.savefig("model_regression_loss.png")

```




```

1 # training metrics
2 scores = model.evaluate(seq_array, label_array, verbose=1, batch_size=bs)
3 print('\nMAE: {}'.format(scores[1]))
4 print('\nR^2: {}'.format(scores[2]))
5
6 y_pred = model.predict(seq_array, verbose=1, batch_size=bs)
7 y_true = label_array
8
9 test_set = pd.DataFrame(y_pred)
10 test_set.to_csv('submit_train.csv', index=None)

```

37/37 [=====] - 7s 200ms/step - loss: 434.2534 - mae: 13.1452 - r2_keras: 0.8532

MAE: 13.145186424255371

R^2: 0.8532044291496277

37/37 [=====] - 8s 199ms/step

MODULE-4: TESTING THE MODEL

We finally apply the Training dataset whose ground truth Remaining useful lifetime is given in a separate dataset into our trained and optimized LSTM model, we obtain the output for each case of aircraft engine based on the aircraft id. The RUL predicted and actual ground truth is compared to analyze the result.

```

1 # Test data validation
2
3 # We pick the last sequence for each id in the test data
4 seq_array_test_last = [test_df[test_df['id'] == id][sequence_cols].values[-sequence_length:]
5 | | | | | | | | | | for id in test_df['id'].unique() if len(test_df[test_df['id'] == id]) >= sequence_length]
6
7 seq_array_test_last = np.asarray(seq_array_test_last).astype(np.float32)
8 print("seq_array_test_last")
9 # print(seq_array_test_last)
10 print(seq_array_test_last.shape)
11
12 # Similarly, we pick the labels
13 # print("y_mask")
14 y_mask = [len(test_df[test_df['id'] == id]) >= sequence_length for id in test_df['id'].unique()]
15 label_array_test_last = test_df.groupby('id')['RUL'].nth(-1)[y_mask].values
16 label_array_test_last = label_array_test_last.reshape(label_array_test_last.shape[0], 1).astype(np.float32)
17 # print(label_array_test_last.shape)
18
19 # if best iteration's model was saved then load and use it
20 if os.path.isfile(model_path):
21     estimator = load_model(model_path, custom_objects={'r2_keras': r2_keras})
22
23 # test metrics
24 scores_test = estimator.evaluate(seq_array_test_last, label_array_test_last, verbose=2)
25 print('\nMAE: {}'.format(scores_test[1]))
26 print('\nR^2: {}'.format(scores_test[2]))
27
28 y_pred_test = estimator.predict(seq_array_test_last)
29 y_true_test = label_array_test_last
30
31
32 pd.set_option('display.max_rows', 1000)
33 test_print = pd.DataFrame()
34 test_print['y_pred'] = y_pred_test.flatten()
35 test_print['y_truth'] = y_true_test.flatten()
36 test_print['diff'] = abs(y_pred_test.flatten() - y_true_test.flatten())
37 test_print['diff(%)'] = abs(y_pred_test.flatten() - y_true_test.flatten())/y_true_test.flatten()
38 print(test_print)

```

```

41 test_set = pd.DataFrame(y_pred_test)
42 test_set.to_csv('submit_test.csv', index=None)
43
44 # Plot in blue color the predicted data and in green color the
45 # actual data to verify visually the accuracy of the model.
46 fig_verify = plt.figure(figsize=(12, 6))
47 plt.plot(y_pred_test, color="blue")
48 plt.plot(y_true_test, color="green")
49 plt.title('prediction')
50 plt.ylabel('value')
51 plt.xlabel('row')
52 plt.legend(['predicted', 'actual data'], loc='upper left')
53 plt.show()
54 fig_verify.savefig("model_regression_verify.png")

```

```


seq_array_test_last
(88, 60, 25)
3/3 - 2s - loss: 353.7944 - mae: 12.7856 - r2_keras: 0.7868 - 2s/epoch - 647ms/step

```

MAE: 12.7855863571167

R²: 0.7868390679359436

	y_pred	y_truth	diff	diff(%)
0	63.190376	69.0	5.809624	0.084197
1	93.219536	82.0	11.219536	0.136824
2	100.930458	91.0	9.930458	0.109126
3	135.100723	93.0	42.100723	0.452696
4	107.485336	91.0	16.485336	0.181158
5	107.071800	95.0	12.071800	0.127072
6	86.003746	96.0	9.996254	0.104128
7	107.637825	97.0	10.637825	0.109668
8	99.441872	124.0	24.558128	0.198049
9	84.269844	95.0	10.730156	0.112949
10	116.333488	83.0	33.333488	0.401608



PERFORMANCE MEASURES

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=0}^n (\hat{y}_i - y_i)^2}$$

where \hat{y}_i and y_i are respectively the value predicted by the model and the ground truth provided by the test set of the dataset of the generic i -th input sample, and n is the number of samples of the test set.

REFERENCES

- [1] On the use of LSTM networks for Predictive maintenance, in smart industries." - 2019 IEEE International Conference of Smart Computing.
- [2] Machine Learning for Equipment failure Prediction and Predictive Maintenance (pm] - shad Griffon - Medium.com
- [3] Towards Data Science - Predictive maintenance of turbofan engines - by Keen Peters
- [4] Predicting the maintenance of aircraft engines using LSTM - International Journal trend in Scientific research and development. ISSN: 2456-6470
- [5] LSTMs Explained: A Complete, Technically Accurate, Conceptual Guide with Keras – Ryan – Medium.com
- [6] Recurrent Neural Networks (RNN) and Long Short-Term Memory (LSTM) – Brandon Rohrer
- [7] Towards Data Science - Illustrated Guide to LSTM's and GRU's: A step by step explanation